

Yuanfan Chen

646-963-8426 | yc2942@cornell.edu | [github](#) | [linkedin](#) | [scholar](#)

EDUCATION

Cornell Tech, Cornell University

Master of Engineering in Computer Science (GPA: 3.9)

New York, NY

Aug. 2025 – May. 2026

University of Toronto

Honours Bachelor of Science, Computer Science Specialist (GPA: 3.9)

Toronto, Canada

Sep. 2020 – May. 2024

PUBLICATIONS

- **Beyond Prediction: Tail-Aware Scheduling for LLM Inference.** Yueying Li, **Yuanfan Chen**, Jiayang Chen, Esha Choukse, Haoran Qiu, G. Edward Suh, Rodrigo Fonseca, Ziv Scully, Udit Gupta. **ICML 2026**. arXiv:2606.18431.
- **Making Sense of DPU Performance for Cloud Data Processing.** Jiasheng Hu, Chihan Cui, **Yuanfan Chen**, Philip A. Bernstein, Jialin Li, Qizhen Zhang. **SoCC 2026** (17th ACM Symposium on Cloud Computing). arXiv:2504.05536.

OPEN SOURCE

vLLM | *ROCm, fp8 MoE*: Added `gelu_tanh` to the AITER fp8 fused-MoE backend; fixed uninitialized padded expert weights that corrupted outputs (#55251).

LMCache | *Observability*: L1/L2 cache-hit attribution and Prometheus counters for the multi-process lookup path (#4734, #4962).

SGLang-Omni | *CUDA IPC, MPS*: Same-GPU weight sharing across data-parallel replicas (#1124); co-developed Restage placement ranking (#2134).

EXPERIENCE

Machine Learning Infra Engineer

Aug 2026 – Present

Character.AI: Inference Optimization

Bay Area, CA

- Own production inference optimization for large-scale **multi-turn, long-context** chat on a **vLLM + LMCache** stack serving the **Gemma**, **DeepSeek**, and **GLM** model families across **thousands of GPUs**, including latest-generation **AMD Instinct** clusters; work spans **KV-cache offloading** and **kernel optimization**.
- Deployed a **hierarchical, distributed KV cache** (GPU → L1 → shared remote tiers) with **LMCache**, raising the **prefix-cache hit rate from 70% to 90%** and eliminating most long-context prefill recompute.
- Drive kernel work from measurement: profile compute- vs. memory-bound phases and per-layer time (attention vs. FFN) against **MFU** ceilings; sustained **~3K prefill tokens/s per MI325X** on a 30B-class Gemma model.
- Built a **real-traffic replay benchmark** to validate every optimization on production workload patterns before rollout.

Machine Learning System Researcher – S4AI Lab

Sep 2025 – Present

Cornell University

New York, NY

- Engineered a scheduling framework integrating **RCPSP modeling**, **EDF**, and **Greedy** to maximize throughput under **multi-class SLO constraints**; derived **stepwise performance functions** via GPU profiling and validated performance by reproducing SOTA baselines (**Slo-serve**, **Sarathi**).
- Leveraged **γ -Boost Queuing Theory** to analyze **light-tailed workloads**, optimizing resource scheduling for prefill versus decode phases; achieved a **300% tail latency improvement** while maintaining Avg E2E and TTFT parity with **SJF (Shortest Job First)**.
- Developed a **Vidur-based simulator** for Multiclass SLO evaluation, conducting a class-level comparative analysis of **Colocation vs. Separation** global scheduling strategies and reproducing **PolyServe** as a baseline.

Machine Learning Engineer

May 2025 – Sep 2025

Tencent: Distributed Training and Inference

Beijing, CN

- Mastered parallelization strategies including **Tensor Parallelism (TP)** and **Data Parallelism (DP)**, and optimized inference engines using **PagedAttention** and **Continuous Batching**.
- Optimized **GQA** and **DSA** kernels on **NVIDIA H200 and B200** GPUs to accelerate LLM inference; profiled memory hierarchies and warp stalls using **Nsight Compute (NCU)** to alleviate bandwidth bottlenecks and maximize hardware utilization.
- Reduced pipeline bubbles by implementing **Chunk Prefill** and batched decode strategies, significantly improving throughput for long-context inference workloads.

Undergraduate Student Researcher – Far Data Lab

Jan 2024 – Feb 2025

University of Toronto

Toronto, CA

- Co-designed **dpBento**, a modular benchmark framework that standardizes performance evaluation across heterogeneous DPUs (**BlueField-2/3**, **OCTEON TX2**), from micro-level primitives to end-to-end cloud database offloading.
- Profiled **RocksDB** compaction and decompression paths end-to-end; offloaded compute-heavy stages to DPU, reducing **host CPU utilization by ~40%** and achieving **~3× decompression throughput** via hardware acceleration.

ACADEMIC SERVICE

- **Invited Reviewer**: IEEE Journal on Selected Areas in Information Theory (JSAIT) – Special Issue on Energy and Data Efficiency in AI (2026). Reviewed submissions focusing on cutting-edge LLM inference systems.